Coevol

Correlated evolution of substitution rates and quantitative traits

Nicolas Lartillot, Raphael Poujol

nicolas.lartillot@univ-lyon1.fr

Version 1.6. June 6, 2021.

Contents

1 Introduction

DNA sequences evolve at a different rate in different species. More generally, several aspects of the substitution process may be subject to variation between organisms: the nucleotide composition (and in particular the GC content), the ratio of non-synonymous over synonymous substitution rates (dN/dS), the ratio of transition over transversion rates, etc. Understanding the causes of such variation is a fundamental question in evolutionary genetics, still open in many respects.

The comparative method represents a natural approach to investigating these issues. Correlations between the substitution rate, dN/dS or GC content on one hand, and body mass, generation time, metabolic rate, or genome size on the other hand, would certainly help discriminating between alternative hypotheses about the causes of rate and substitution pattern variation. In all cases, correlations should be corrected for non-independence due to phylogenetic inertia, as would be done in any other phylogenetic correlation study.

Classical comparative methodologies, however, only apply to characters that can be directly observed and measured: body-size, life-history traits, or any other morphological or biological trait. Here in contrast, we need to make correlations between observable traits (e.g. body size) and parameters of the substitution process (such as the substitution rate, or the mutation bias, or dN/dS) which manifest themselves only indirectly via the DNA sequences of the alignment.

The aim of Coevol is to propose an integrated statistical framework for investigating these kinds of comparative questions, using a combination of DNA sequences, comparative data and paleontological knowledge.

Coevol: the comparative method applied to molecular data

Coevol is a Bayesian inference program using Markov Chain Monte Carlo methods. It can be seen as a fusion between classical phylogenetic models for nucleotides or codons (Felsenstein, 1981), autocorrelated relaxed clocks for molecular dating (Thorne et al., 1998) and comparative models based on the idea of Brownian processes and phylogenetically independent contrasts (Felsenstein, 1985; Martins and Hansen, 1997; Harvey and Pagel, 1991).

In Coevol, the estimation works by conditioning a probabilistic model simultaneously on a sequence alignment, a matrix of quantitative characters (such as morphological data or life-history traits) and fossil calibrations. The model assumes correlated evolution of the rate of substitution, or other parameters of the substitution process such as GC content or dN/dS, and the quantitative traits, all of which are jointly modeled as a multivariate Brownian process. The program then estimates the correlation structure between these variables (i.e. the covariance matrix of the Brownian process) and simultaneously reconstructs divergence

times and ancestral traits along the phylogeny.

The most straightforward application of Coevol is probably the identification of quantitative traits correlating with variation in substitution rates or substitution patterns. However, Coevol can also be used for reconstructing ancestral traits (like BayesTrait, Pagel, 1999), or for estimating divergence times (like Beast or Multidivtime, Drummond and Rambaut, 2007; Thorne et al., 1998). Thanks to the integrative approach adopted here, estimation of any of these components of the model (correlations, divergence times, or ancestral traits) will automatically integrate uncertainty, as well as potentially relevant information, coming from the other components.

The method has been introduced in Lartillot and Poujol (2011) and has been used, among other things, for identifying the life-history correlates of dN/dS and GC content in placental nuclear genomes (Lartillot, 2013) or for estimating divergence times and ancestral body masses in mammals (Lartillot and Delsuc, 2012).

TipCoevol: A tip-dating version of the molecular comparative method

The models implemented in the current version of `coevol` proceed from a node-dating approach, in which fossil calibrations are part of the prior and are imposed on internal nodes of an ultrametric tree containing only extant taxa (Thorne and Kishino, 2002; Yang and Rannala, 2006; Drummond et al., 2006). Recently, tip-dating, or total-evidence dating, has emerged as an alternative molecular dating paradigm, which explicitly accounts for the fact that fossils and extant taxa are part of the same macro-evolutionary process (Pyron, 2010; Ronquist et al., 2012; Stadler and Yang, 2013; Heath et al., 2014). This macro-evolutionary process is typically formalized in terms of a birth-death process with serial species sampling through time (Stadler, 2010; Stadler and Yang, 2013; Heath et al., 2014). In the context of tip-dating, fossils are represented as tips of the phylogenetic tree, whose age is constrained to be within some time interval determined based on paleontological evidence. The fossilized birth-death (FBD) model is a more general and more accurate variant of this approach, allowing fossils to be direct ancestors to other fossils or to extant taxa (Heath et al., 2014). In both cases (tip-dating or FBD), calibration of divergence times by fossils is implicit, being mediated by the prior on divergence times induced by the serial or fossilized birth-death process (Heath et al., 2014) and, in the case where discrete morphological characters are used, the relaxed morphological clock (Pyron, 2010; Ronquist et al., 2012). For those theoretical reasons, tip-dating is currently emerging as an attractive alternative to node-dating.

A tip-dating version of `coevol` is provided as a separate program, named `tipcoevol`. This program also implements the mixed relaxed molecular clock introduced in Lartillot et al. (2016). According to this relaxed clock model, rate variation has both a long-term Brownian behavior (presumably driven by long-term changes in life-history strategies) and

short-term variations, the latter being captured by a whitenoise process. By estimating the relative variance of its two components directly from the sequence alignment, the mixed clock interpolates between the uncorrelated and correlated relaxed clocks, containing them as particular cases. Applying this mixed clock model to an empirical dataset leads to a simple measure of how the total variation in substitution rate across the tree is partitioned into the Brownian and the whitenoise components, which then yields potentially useful insight about how rate variation unfolds over different timescales.

Note that, in practice, if most of the variance is contributed by the Brownian or by the white-noise component, this means that the program has effectively selected the relevant pure clock automatically. In other words, no explicit model comparison is necessary here, to select among alternative, autocorrelated or uncorrelated clocks.

In `tipcoevol`, the mixed relaxed clock is implemented in the context of a multivariate Brownian process. In this context, the Brownian component of the mixed clock will be one of the components of a multivariate process, while the other components of the process will map to whichever quantitative traits are used in the analysis (such as life-history traits, as in `coevol`). If used without quantitative traits, the `tipcoevol` program will reduce to the univariate mixed-clock model introduced in Lartillot et al. (2016).

Ancov: Ancestral state reconstruction by comparative regression

A comparative program independent of Coevol, called `ancov` (ancestral covariance), was introduced in Lartillot (2014). The aim of this program is to reconstruct ancestral traits based on a Gaussian regression model implemented using a Kalman filtering algorithm. A possible application, such as illustrated in the article, is the reconstruction of ancestral growth temperatures in prokaryotes based on inferred ancestral proteome or rRNA composition, both of which are known to correlate very strongly with growth temperature (Groussin and Gouy, 2011).

Although this question would ideally require a one-step integrative modeling, such as `coevol`, the idea of `ancov` is to proceed in two steps: first estimate ancestral compositions using a phylogenetic methods, and second, rely on the trait / composition correlation in extant species, as well as on inferred ancestral compositions, to reconstruct ancestral temperatures. Thus, Ancov can be seen as an approximate method, but one that still corrects for phylogenetic inertia for reconstructing ancestral traits based on ancestral compositions.

More generally, the program can be used for any type of ancestral reconstruction using a comparative regression approach, such as already formalized in other previously published programs and applications (Martins and Hansen, 1997; Pagel, 1999; Organ et al., 2007; Franks et al., 2012).

Details about the `ancov` program are given below, in section 7.

2 General features and some practical considerations

2.1 Sequence data: codons or nucleotides

Coevol can run on nucleotide and codon alignments. In the case of nucleotides, the model allows for variation in the overall substitution rate or, using specific options, for variation in the ratio between transition and transversion rates (κ) or in the equilibrium GC content. When applied to codon data, the model will also reconstruct phylogenetic variation in synonymous and non-synonymous rates (dS and dN) or in their ratio $\omega = dN/dS$. Finally, it is possible to apply the program to amino acid recoded sequences, so as to estimate variation of the radical over conservative amino-acid replacement rates (Nabholz et al., 2013).

2.2 Fixing or estimating divergence times

Estimation can be conducted under a fixed time-calibrated phylogeny or, alternatively, divergence times can be co-estimated. Fixing divergence times may inflate significance of some of the correlations (Diaz-Uriarte and Garland, 1998). Conversely, divergence time estimation can sometimes be a limiting factor for MCMC convergence, in particular for large trees with a large number of (potentially conflicting) calibrations. Thus, it can sometimes be useful, at least in a first step, to constrain the time-calibrated phylogeny.

When divergence times are not a priori constrained, fossil calibrations can be specified, although they are optional. If no calibration is specified, the tree will be dated relative to the age of the root (by convention, time scale is defined so that the age of the root is equal to 1). However, one should keep in mind that relative dating without calibration may not be so reliable, especially for deep phylogenies (e.g. at the level of orders). If your interest is in correlations between the absolute substitution rate and quantitative traits, relative dating may give you poor estimates of divergence times, and therefore, correlatively, poor estimates of ancestral substitution rates (rates and times are two sides of the same coin: only their product is directly measurable). In such cases, specifying calibrations for constraining divergence time estimation is probably a good thing, and checking the sensitivity of the analysis to fossil calibrations is even better (Lartillot and Delsuc, 2012).

On the other hand, if your interest is in correlating intensive substitution variables (ratios between rates, such as the ratio between transition and transversion rates, or dN/dS, or GC content) with quantitative traits, then divergence time estimation is less critical. It still has an influence in principle (because the Brownian motion, as a diffusion process, is defined as a function of evolutionary time) but, on the other hand, the errors on rate estimates directly induced by the errors on divergence times will cancel out in the ratios of rates.

2.3 Quantitative traits

Quantitative characters will be automatically log-transformed by the program. That is, the program will assume that the natural logarithm of the characters given as an input will evolve according to a Brownian motion. The initial justification of this automatic log-transformation is that, for characters such as body size, life-history traits, or substitution rates, it is reasonable to assume allometric (log-linear) relations. However, this means that, if you want to conduct correlation studies with other types of characters, you may need to apply some transformation to them before giving them to coevol. For instance, characters which should evolve as Brownian motions on a linear scale should be exponentiated (so as to offset the log transformation applied by coevol). As another example, for characters such as fractions, between 0 and 1 (strictly), it is natural to assume a log-it transformation. Thus, if x is your character, then one would define

$$y = \ln \frac{x}{1-x}$$

and assume Brownian evolution of y. In that case, what you should give coevol is not y but z defined as:

$$z = \frac{x}{1-x}.$$

Coevol will automatically take the natural logarithm of z, thus effectively working with y.

2.4 How long should it run?

It is generally difficult to know beforehand how long a chain should run. In the case of coevol, our general experience is that it takes one hundred to a few hundred points to reach convergence, and a sample of 1 000 points (after burnin) generally gives reasonable qualitative estimates. High-quality 'publication-grade' runs, on the other hand, may require longer runs. With 100 species and 5 000 aligned positions, this means roughly an overnight analysis for having a good qualitative idea and several days to one or two weeks for high-quality results. However, these are just orders of magnitude. Different datasets, or different models, may require different numbers of cycles before reaching convergence and may display very different mixing behaviors. The best is to rely on more objective measures, such as effective sample size and reproducibility of the results across independent runs started from random initial conditions (all of which are explained in section 4).

For a faster 'research cycle', it is often a good idea to first domesticate the problem at hand with smaller datasets. In particular, if your dataset contains a large number of taxa (several hundreds), it can be a good thing to first run the program on a subsample of ~ 100

taxa. Less critically, using fixed time-calibrated phylogenies can help in a first step, before turning to more ambitious models and analyses. More generally, one should not hesitate to first capture the essential points using fast qualitative runs (along the lines indicated above), although *always* running two independent chains in parallel and *always* visualizing the tracefiles to check convergence. Coevol allows you to stop and restart chains at will, thus it is easy to make quick runs under several conditions, stop them when you get a good qualitative idea of the outcome, and finally, once the broad lines of an interesting series of results appear to be within reach, elongate the most interesting chains for final high-quality results.

3 Input data format

3.1 Sequences

The format recognized by Coevol is a generalization of the PHYLIP format:

```
<number_of_taxa> <number_of_sites>
taxon1 sequence1...
taxon2 sequence2...
...
```

Taxon names may contain more than 10 characters. Sequences can be interrupted by space and tab, but not by return characters. Sequences can be interleaved, in which case the taxon names may or may not be repeated in each block.

The following characters will all be considered equivalent to missing data: "-", "?", "$", ".", "*", "X", "x", as well as the degenerate bases of nucleic acid sequences ("B", "D", "H", "K", "M", "N", "R", "S", "V", "W", "Y"). Upper or lower case sequences are both recognized, and the case matters for taxon names (but not for nucleotides). Codon alignments should be provided as simple nucleotide alignments, with the only constraint that the number of positions should be a multiple of 3. The `<number_of_sites>` field should be the number of aligned nucleotide sequences (not the number of codon positions). Stop codons are allowed, as well as incomplete codons (e.g. 'CA-' or 'CA?'), but will be replaced by unknown states.

3.2 Trees

Trees should be in the Newick format, and should be rooted, e.g.:

```
((taxon1:0.1,(taxon2,taxon3):2.3):1.2,(taxon4:0.5,taxon5:0.2):0.23);
```

Branch lengths can be specified. If the tree is ultrametric (if it is a time-calibrated phylogeny), the branch lengths specified in the tree file will be used as initial divergence times (or, using the -fixbl option, can be fixed as a prior constraint). Trees should always be followed by a ';'. Taxon names should correspond to the names specified in the data matrix (case sensitive). If some names are present in the tree, but not in the matrix, the corresponding taxa will be pruned out of the tree. That is, the spanning subtree containing all the taxa mentioned in the data matrix will be considered as the input tree. Conversely, if some taxa are present in the data matrix, but not in the input tree, the program will exit with an error message.

In the case of tip dating, the tree is not ultrametric, since fossils are tips, whose age is thus strictly positive. The tree given to the program should specify branch lengths and should be compatible with the intervals for the fossil ages given by the calibration file.

3.3 Matrix of characters

The list of continuous characters should be in a separate file, formatted as follows :

```
#TRAITS
<number_of_taxa> <number_of_characters> NAME1 NAME2 ...
taxon1 character_1 character_2 ...
taxon2 character_1 character_2 ...
...
```

Thus, each trait should be given a name (e.g. `NAME1` = longevity, `NAME2` = body_mass), names which will be used by coevol in the display of the correlation matrices. All characters should be strictly positive (as mentioned above, they will be automatically log-transformed). Missing data can be indicated by setting the character at -1.

The `#TRAITS` keyword is present for historical reasons. In the first version of the program, no name was indicated and the format (still recognized by this version) was more simply:

```
<number_of_taxa> <number_of_characters>
taxon1 character_1 character_2 ...
taxon2 character_1 character_2 ...
...
```

Both formats are recognized, but the important thing is to put the `#TRAITS` keyword as a header to the file whenever traits are given explicit names.

3.4 Calibrations

Calibrations should be given in a separate file. The format is such that specific internal nodes of the phylogeny are indicated by giving the names of two terminal taxa that have this node as their last common ancestor – as follows:

```
<ncalib>
<taxon1a> <taxon1b> <upper_limit> <lower_limit>
<taxon2a> <taxon2b> <upper_limit> <lower_limit>
...
```

Upper or lower limits can be set equal to -1, in which case no limit is enforced. For example:

```
taxon1 taxon2 70 50
```

specifies an upper and a lower constraint, thus the interval [50,70] My,

```
taxon1 taxon2 -1 50
```

means that the node of the last commmon ancestor of taxon1 and taxon2 should be older than 50 Million years (My).

```
taxon1 taxon2 70 -1
```

only specifies an upper constraint of 70 My, and no lower constraint.

The root can be calibrated by specifying an upper and/or a lower bound for the root node in the calibration file, exactly like for other nodes. In addition, a gamma prior for the age of the root can be imposed, by specifying a mean and a standard deviation for the age of the root in the command for launching the program (see below). The two priors are not mutually exclusive (in which case the gamma distribution will be truncated according to the specified hard bounds). The absolute age of the root is often the most challenging parameter of the model, in terms of MCMC convergence and mixing.

In the case of tip-dating, only tips should be calibrated. This can be done using a specialization of the format shown above: since a tip is the most recent common ancestor of itself, it can be referred to by giving its name twice in the calibration file. Thus, the calibration file should list all fossils included in the analysis, and for each of them, give an upper and a lower limit for its age:

```
<ncalib>
<tip_taxon1> <tip_taxon1> <upper_limit> <lower_limit>
<tip_taxon2> <tip_taxon2> <upper_limit> <lower_limit>
...
```

3.5 Ancestral quantitative data (for the `ancov` program)

Ancestral quantitative data, i.e. traits that are known both in extant species and in ancestors corresponding to internal nodes of the phylogeny, should be specified like fossil calibrations: nodes are specified by giving the names of two taxa that have this node as their last common ancestor. A terminal node is specified by giving the name of the corresponding taxon twice. For instance, for a K-dimensional trait specified for a total of N nodes of the phylogeny:

```
<N> <K>
<taxon1a> <taxon1b> <x_11> <x_12> ... <x_1K>
<taxon2a> <taxon2b> <x_21> <x_22> ... <x_2K>
...
```

4 Running a chain: `coevol`

4.1 A running example

Running the `coevol` program will produce a series of points drawn from the posterior distribution over the parameters of the model. Each point defines a detailed model configuration (divergence times, covariance matrix, ancestral traits, etc.). The series of points defines a chain.

An example, analyzed in more details in Lartillot and Delsuc (2012) and Lartillot (2013), is provided with the program. The `coevol/data/plac/` directory contains the following files:

- `plac.ali`: a multiple sequence alignment (17 nuclear genes, 73 placental mammals). This dataset is a codon alignment (amino-acid sequences have been aligned, and this protein alignment has then been used as a guide for aligning the nucleotide sequences without disrupting the reading frame).
- `plac4fold.ali`: The four-fold degenerate third coding positions of the same 17 nuclear genes in 73 placental mammals. This dataset is useful for investigating the patterns of rate (and GC content) variation at putatively neutral positions.
- `plac.tree`: a phylogenetic tree,
- `plac.lht`: a matrix of characters (female age at sexual maturity, adult body mass, and maximum recorder lifespan) obtained from the AnAge database (de Magalhaes and Costa, 2009).
- `plac.calib`: fossil calibrations.

As suggested above, a reasonable starting point will be to first analyze this dataset with a pure nucleotide model. Codon models will be considered in a second step. The command for our first nucleotide analysis is then:

```
coevol -d plac4fold.ali -t plac.tree -fixtimes -c plac.lht plac1
```

The `-d` option is for specifying the dataset, `-t` the tree and `-c` the quantitative character matrix. Here, `plac.tree` is a chronogram (a time-calibrated phylogeny) independently reconstructed using a relaxed clock method implemented in the phylobayes program (Lartillot et al., 2009). We can use this time-calibrated phylogeny as a prior constraint and ask the program to not reestimate divergence times (using the `-fixtimes` command).

A series of files will be produced with a variety of extensions. The most important are:

- `plac1.trace`: the trace file, containing a few relevant summary statistics (e.g. log-likelihood, total length of the tree, age of the root, covariances, etc);

- `plac1.chain`: containing the detailed parameter configurations visited during the run. This file is used by `readcoevol` for computing posterior averages.

Chains will run as long as allowed. They can be interrupted at any time, and then restarted, in which case they will resume from the last check-point (last point saved before the interruption). To soft-stop a chain, just open the `<name>.run` file, and replace the 1 in it by a 0. Under linux, this can be done with the simple following command:

```
echo 0 > plac1.run
```

The chain will finish the current cycle before exiting.

To restart a chain:

```
coevol  plac1
```

Most often, when chains are killed (e.g. because of a time-out on a cluster), they can be restarted in this way. Be careful, however, not to restart an already running chain.

4.2 Alternative models

Coestimation of divergence times using fossil calibrations

For our first chain above, plac1, we have chosen to fix divergence times. Alternatively, we can use the fossil calibrations provided in the `plac.calib` file:

```
coevol -d plac4fold.ali -t plac.tree -cal plac.calib 100 100 -bd -c plac.lht calplac1
```

The option is `-cal`, followed by the name of the file containing the fossil calibrations (see Input data format), followed by two numbers specifying the mean and the standard deviation of the prior over the age of the root (which is a gamma distribution). Here, the root is the last common ancestor of placentals. We opt for a broad prior of mean 100 Myr and standard deviation 100 Myr. In this particular case, the prior is thus an exponential distribution of mean 100. We also opt for a birth-death prior on divergence times (`-bd` option). If no prior were specified, this would have be a uniform prior on divergence times.

In principle, a birth-death has a better justification than a uniform in terms of the underlying diversification process (although assuming constant extinction and speciation rates is probably naive). Uniform priors are paradoxically more informative and will often tend to make more compact trees, with a more recent root. On the other hand, MCMC convergence, in particular on the age of the root tends to be faster the uniform prior.

GC content

There are interesting questions about GC content variation in placentals. To investigate them, we can allow for variation in equilibrium GC frequency (GC^*) along the phylogeny using the `-gc` command:

```
coevol -d plac4fold.ali -t plac.tree -fixtimes -c plac.lht -gc placgc1
```

Codon model: dN/dS

Another interesting option in the present case is the codon model: correlating, not just the total substitution rate, but separately the synonymous substitution rate dS and the ratio of the non-synonymous over the synonymous substitution rates $\omega = dN/dS$ with body size and life-history traits. To do this, we use the `-dsom` ("dS and omega") option:

```
coevol -d plac.ali -t plac.tree -fixtimes -c plac.lht -dsom placdsom1
```

Is is possible to combine dS, dN/dS and GC^* in the same analysis:

```
coevol -d plac.ali -t plac.tree -fixtimes -c plac.lht -dsom -gc placdsomgc1
```

Whichever model configuration is used, runs should be duplicated (`calplac2`, `placgc2`, `placdsom2`, etc), and should be run for sufficiently long (see above). A good starting point is to let the chains run overnight before having a first look at what they give.

4.3 Checking convergence

The MCMC sampler saves one point after each cycle. A cycle itself represents a set of complex and integrated series of updates of the parameters of the model (divergence times, covariance matrix, ancestral traits and rates, etc).

Convergence can first be visually assessed by plotting the summary statistics recorded in the trace file as a function of the number of iterations. This can be done using simple linux utilities, such as `gnuplot`. Alternatively, the trace file of coevol is compatible with the `Tracer` program of the Beast software (Drummond and Rambaut, 2007). You can therefore use `Tracer` to check convergence and estimate burn-in and effective sample size (`tracecomp`, introduced below, does similar things but has a more primitive interface).

The following statistics, tabulated in the .trace files, are particularly relevant for checking convergence:

- log prior (column 1)
- log likelihood (column 2, `lnL`; in general, converges relatively quickly)

- tree length (column 3; for codon models, this is the synonymous length, i.e. the expected number of synonymous substitutions per site across the whole tree)
- for codon models, mean dN/dS or total dN over the tree
- at least checking a few entries of the covariance matrix (the Σ_{ij} in the trace file).
- the age of the root (in those cases where divergence times are also estimated).

Plotting trace files for the placental mammals example

Plotting these statistics for the `calplac1` and `calplac2` runs, after one night (630 points), would give something like on figure 1. Convergence appears to be very rapid in the present case (burn-in of less than 100 points), except for the age of the root, although the main problem with this root age is not so much convergence than mixing: low frequency oscillations are visible on the trace plot, suggesting that it may take more cycles for decorrelating this root age.

After one night, the codon model (with dN/dS and GC^*) has produced 400 points. The essential qualitative aspects of the correlation structure between dS, dN/dS and GC^* and life-history traits (such as estimated using `readcoevol`, introduced in the next section) can already be seen based on the last 300 points of the run. High-quality results will be obtained after running the program for approximately 1 week.

The `tracecomp` program

In addition to visual checks, a more quantitative assessment of convergence and mixing can be performed using the `tracecomp` program:

```
tracecomp -x 100  calplac1 calplac2
```

`tracecomp` will produce an output summarizing the discrepancies and the effective sizes estimated for each column of the trace file. The discrepancy d is defined as

$$d = 2|\mu_1 - \mu_2|/(\sigma_1 + \sigma_2),$$

where μ_2 and μ_2 are the means and σ_1 and σ_2 the standard deviations associated with a particular column, for the two chains. The effective size is evaluated using the method of Geyer (1992). Some guidelines for evaluating the quality of the samples:

- maxdiff < 0.1 and minimum effective size > 300: good run;

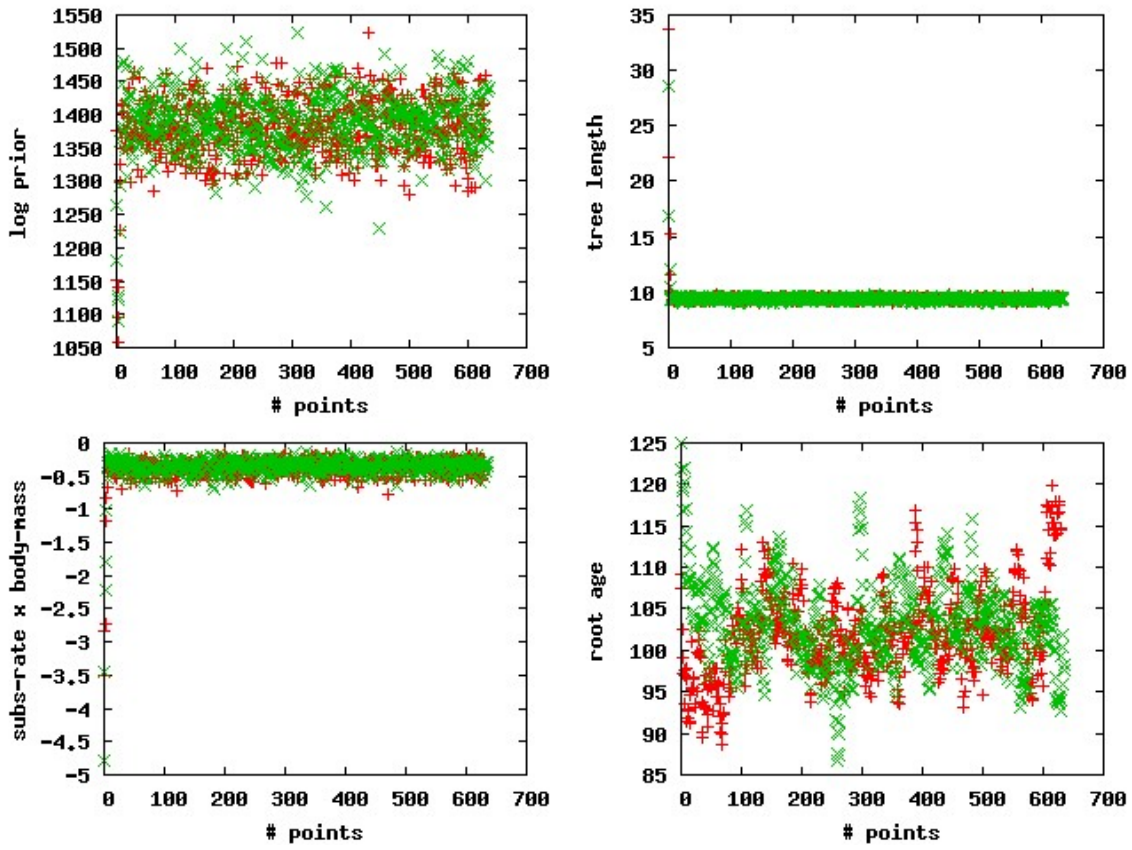

Figure 1. Traceplots for the `placdsgc1` chain showing, as a function the number of points saved in the tracefile, the log prior, total tree length, covariance between dS and body mass, and the age of the root

- maxdiff < 0.3 and minimum effective size > 50: acceptable run (yielding qualitatively correct results).

Often, most summary statistics of the trace file will have large effective sample sizes, except one or two of them (in particular, the age of the root).

`tracecomp` applied to our placental mammals example

In the present case (`calplac1`, after one night and 630 points), the output of `tracecomp` reads like:

```
setting upper limit to : 630
burnin : 100
stop : 630
name                effsize     rel_diff

#logprior           406         0.164942
lnL                 449         0.25871
rate                514         0.240122
sigma_0_0_1         530         0.0710222
sigma_0_0_2         530         0.209856
sigma_0_0_3         530         0.19284
...
sigma_0_3_3         530         0.134073
rootage             87          0.0615232
...
```

suggesting that, already after one night, we get relatively good qualitative results: relative discrepancies all below 0.26 and estimated effective sample sizes close to true sample size (530), except for the age of the root. After 4 days, we have a chain of approximately 5200 points, with excellent convergence and mixing diagnostics (effective sample size larger than 500 for the age of the root and 1000 for all other statistics, discrepancies smaller than 0.02).

5 Post-analysis (readcoevol)

5.1 Obtaining posterior averages

Once a chain has been obtained, and the burnin has been determined, correlations, divergence times and ancestral traits can be estimated using readcoevol:

```
readcoevol -x <burn-in> [<every> <until>] <chain_name>
```

By default, `<burn-in>` is equal to one tenth of the total size of the chain (number of points saved), `<every>` = 1 and `<until>` is equal to the size of the chain. However, it is preferable to always specify the burnin. Thus, for instance:

```
-x 300
```

defines a burn-in of 300, computing averages based on all of the remaining points,

```
-x 300 10
```

a burn-in of 300, taking one every 10 points, up to the end of the chain, and

```
-x 300 1 5300
```

a burn-in of 300, taking one every point, up to the 5300th point of the chain (or less, if the chain is shorter).

The placental mammal example

Following our example:

```
readcoevol -x 200 plac1
```

will discard the first 200 points and will compute posterior estimates based on the 5000 remaining points. It will produce a few additional files, contaning:

- `plac1.cov`: the estimated covariance matrix, and associated correlation coefficients and posterior probabilities.
- `plac1.postmeandates.tre`: posterior means and confidence intervals for divergence times
- `plac1.postmean1.tre`: ancestral reconstruction for quantitative trait 1 (same thing for trait 2, etc)

5.2 Correlation analysis: the `.cov` file

As a more interesting example, we could now look at the correlations obtained under the model allowing for simultaneous variation in substitution rate and in GC^*. After running `readcoevol` on `placgc1`, the resulting `placgc1.cov` file would read as:

```
entries are in the following order:
dS
gc
maturity
mass
longevity

covariances
  0.492   0.166  -0.454   -1.38  -0.364
  0.166    1.39  -0.127    -1.1 -0.0636
 -0.454  -0.127     1.7    2.66   0.668
  -1.38    -1.1    2.66    12.6    1.96
 -0.364 -0.0636   0.668    1.96   0.588

correlation coefficients
      1   0.202  -0.496  -0.553  -0.675
  0.202       1 -0.0796   -0.26 -0.0688
 -0.496 -0.0796       1   0.573   0.666
 -0.553   -0.26   0.573       1   0.719
 -0.675 -0.0688   0.666   0.719       1

posterior probs
  -     0.9     0.0019    0       0
0.9       -     0.34    0.036   0.31
0.0019  0.34      -       1       1
  0     0.036     1       -       1
  0     0.31      1       1       -
```

The entries of the 5x5 matrices are in the order specified in the header. The posterior probabilities of a positive correlation (pp) are particularly important. A pp close to 1 means a strong statistical support for a positive correlation, and a pp close to 0 a supported negative correlation (the posterior probabilities for a negative correlation are given by $1 - pp$).

In this example, The first most obvious (albeit not surprising) result is the strong correlation between the three life-history traits (all pp indistinguishable from 1). In this respect,

coevol can be seen as a Bayesian equivalent of classical methods based on phylogenetically independent contrast for estimating correlations between quantitative traits.

As for the correlation between substitution variables (rates) and quantitative traits, we can see that dS displays statistically supported negative correlations with all life-history traits (pp close to 0). In addition, GC^* is negatively correlated with body mass ($pp = 0.036$). The correlation of GC^* with body mass could be due either to a systematic variation in mutation bias as a function of body size (more AT biased mutation process in large-bodied mammals). Alternatively, as further discussed in Lartillot (2013), this could be the result of modulations of the intensity of GC-biased gene conversion as a function of population size, itself correlated with body size.

The other potentially interesting information is the correlation coefficient (r): we see that the negative correlation between dS and longevity is strong ($r = -0.67$). In other words, longevity explains $r^2 = 45\%$ of the variation in dS. This is nearly as strong as the correlation between body mass and longevity in absolute value ($r = 0.72$). The correlation between GC^* and body mass, on the other hand, is relatively weak ($r = -0.26$).

5.3 Multiple regression / partial correlations

The covariances, correlations and posterior probabilities discussed in the last subsection are *marginal*, in the sense that each pairwise correlation (say, between dS and body mass) potentially includes indirect and simultaneous correlation of the two variables with a third one (say, longevity). In contrast, we may be interested in the *partial* correlations. There are several ways of defining and computing partial correlations: using the precision matrix, or explicitly controlling for one or a few traits or rates.

Precision matrix and maximally controlled correlations

The precision matrix is the inverse of the covariance matrix: $\Omega = \Sigma^{-1}$ (Dempster, 1972). For each i, j, $i \neq j$, $-\Omega_{ij}$ is equal to the partial covariance between variables i and j, that is, controlling for all other variables represented in the multivariate process. The partial correlation coefficients are given by (Wong et al., 2003):

$$r_{ij} = -\frac{\Omega_{ij}}{\sqrt{\Omega_{ii}\Omega_{jj}}}.$$

Because all other variables are being controlled, we may call these correlations *maximally controlled correlations*.

The maximally controlled correlations are given after the marginal correlations in the cov file. In our example:

```
precisions
   4.99  -0.421   0.187  0.0597    2.63
 -0.421    1.01 -0.0287   0.139    -0.6
  0.187 -0.0287    1.28 -0.0892   -1.04
 0.0597   0.139 -0.0892   0.218  -0.584
   2.63    -0.6   -1.04  -0.584    6.64

partial correlation coefficients
     -1    0.18 -0.0684  -0.056  -0.448
   0.18      -1  0.0223  -0.289   0.224
-0.0684  0.0223      -1   0.168   0.356
 -0.056  -0.289   0.168      -1    0.48
 -0.448   0.224   0.356    0.48      -1

posterior probs
  -     0.85    0.34    0.39    0.0094
0.85      -     0.54    0.032   0.91
0.34    0.54      -     0.9     0.99
0.39    0.032   0.9       -       1
0.0094  0.91    0.99      1       -
```

Here, we see that the only partial correlation coefficient between dS and life-history traits that has some statistical support is the one between dS and longevity ($pp = 0.0094$, $r = -0.448$). The marginal correlation we have seen between dS and body mass or maturity is therefore most probably indirect, mediated by the correlation of longevity with both dS and body mass. The correlation between GC^* and body mass is robust to controlling for other traits and for dS ($pp = 0.032$).

Custom partial correlations: the `-partial` option and the `.controlcov` file

Partial correlation coefficients obtained using the precision matrix represent a very convenient approach to multiple regression. However, in some cases, this method may not be optimal. Controlling for too many parameters sometimes results in a loss of power. In addition, from a logical point of view, exactly what needs to be controlled for depends on the specific question being asked.

Here, for instance, we may want to more specifically test whether the correlation between dS and body mass is entirely mediated by the indirect correlation of both variables with longevity. Thus, we would like to control *only* for longevity. Such customized multiple regressions can be done with the help of the `-partial` option. After the `-partial` keyword,

one should specify, among all the variables of the analysis, which should be included in the correlation (1) and which should be controlled for (0). In our example, the entries of the process are dS, GC^*, maturity, body mass and longevity. If we want to control only for longevity, which is entry number 5, this would be encoded by 11110.

```
readcoevol -x 200 1 -partial 11110 placdsgc1
```

This will produce a .controlcov file containing the correlations controlled for longevity and the associated posterior probabilities. The matrix is now 4x4, since the entry corresponding to longevity is now missing:

```
entries are in the following order:
dS
gc
maturity
mass

reduced correl

covariances
  0.217  0.0453  -0.027  -0.181
 0.0453   0.987  0.0271  -0.737
 -0.027  0.0271   0.852   0.348
 -0.181  -0.737   0.348    5.34

correlation coefficients
      1   0.102 -0.0637  -0.164
  0.102       1  0.0308  -0.318
-0.0637  0.0308       1   0.161
 -0.164  -0.318   0.161       1

posterior probs
  -     0.69    0.35    0.2
0.69      -     0.57    0.025
0.35    0.57      -     0.89
0.2     0.025   0.89      -
```

From this .controlcov file, we can see that controlling only for longevity indeed results in a complete loss of significance of the correlation between dS and body mass ($pp = 0.2$).

It is possible to control for several variables. In the present case, we can check that, conversely, controlling for body mass and maturity does not compromise the strength and the statistical support of the correlation observed beween dS and longevity:

```
readcoevol -x 300 1 -partial 11001 placdsgc1
```

Here, the command will result in a `placdsgc1.controlcov` file containing a 3x3 matrix:

```
entries are in the following order:
dS
gc
longevity

reduced correl

covariances
  0.252 -0.003    -0.0987
-0.0031  0.854    0.0451
-0.0987  0.0451   0.199

correlation coefficients
      1 -0.0012    -0.441
-0.0012         1   0.109
 -0.441   0.109       1

posterior probs
  -     0.49    0.0075
0.49      -     0.72
0.0075  0.72      -
```

The correlation between dS and longevity indeed remains statistically supported upon controlling for maturity and body mass ($pp = 0.0075$). Which we already knew, since we have already seen above that the maximally controlled correlation between dS and longevity (thus, also controlled for GC^* in addition to maturity and body mass) is supported.

The `.marginalcov` file

The `-partial` option produces yet another file, with the `.marginalcov` extension. This file contains information about the marginal covariance matrix obtained by deleting the entries marked `0`. Thus, for instance:

```
readcoevol -x 200 1 -partial 10111 placdsgc1
```

will result, in addition to what we have just seen (i.e. in a `placdsgc1.controlcov` file in which all correlations will be controlled for GC^*) in a file named `placdsgc1.marginalcov` and contaning the 4x4 matrix of the marginal correlations between dS, maturity, body mass and longevity (not shown here). These latter correlations are *not* controlled for variation

in GC^*. Therefore, this 4x4 matrix is exactly the submatrix obtained from the initial 5x5 matrix Σ by deleting the second raw and second column (which you can check by comparing this `.marginalcov` file and the original `.cov` file). On the other hand, the ensuing precisions and partial correlation coefficients are now partial only among dS and the three quantitative traits. In other words, the partial correlation matrix of dS and longevity displayed in this `.marginalcov` file controls for maturity and body mass, but does not control for GC^*.

General rules for partial correlations / multiple regressions

To summarize, the general rules are as follows:

- the marginal correlations in the `.cov` file do not control for anything;
- the partial correlations in the `.cov` file control for all other traits and rates;
- the entries marked by `0` are the only one controlled for in the `.controlcov` file;
- the entries marked by `0` are the only one *not* controlled for in the `.marginalcov` file.

At first, it may take some thinking to determine exactly what control should be made, depending on the question of interest. In addition, there are obviously redundant ways of obtaining the exact same controlled correlation coefficients here. Altogether, it is perhaps easiest to consider the precision matrix as the most straightforward approach, although it may sometimes result in low power. The `.controlcov` approach is ideal for controlling for a specific set of traits. Finally, the `marginalcov` approach is useful, e.g., for analyzing each aspect of the substitution process (dS, GC^*, dN/dS, etc) in isolation and, in each case, performing a complete multiple regression of the selected variable against all quantitative traits.

5.4 Reconstructing ancestral traits and divergence times

The `readcoevol` program will also produce a series of files containing trees and tabulated lists of ancestral rates, dates and traits. In all case, the posterior summaries can be of several types: mean, standard deviation, or 95% credibility interval, and can be computed on a logarithmic or on a linear scale. Posterior means and standard deviations are perhaps more adequate when computed on the log of the values of interest, whereas the 95% CI can more safely be computed either on a logarithmic or on a linear scale.

By default, the program outputs the 95 % CI on a linear scale, but this can be tuned by using the following options, for activating or deactivating mean or median point estimates, standard deviations, CI, using a logarithmic or a linear scale, or printing the estimated trait values associated to leaf and/or internal nodes:

- `+mean / -mean`
- `+med / -med`
- `+stdev / -stdev`
- `+ci / -ci`
- `+log / -log`
- `+leaf / -leaf`
- `+internal / -internal`

All these options can be freely combined with each other.

Two types of files will be produced, trees (extension `.tre`) or tabulated files (extension `.tab`). The tree files will represent all this information in the newick format. The `.tab` files will tabulate the values for each node of the tree. In the tabulated files, each node of the tree is referred to by specifying two terminal taxa that have this node as their last common ancestor.

For example, assuming that we have obtained 95 % CI on a linear scale, for a data set with 3 taxa, and that the reconstruction of the first quantitative trait included in the character matrix would be something like:

```
            |--------- Taxon1 (40,150)
|-----------| (56,80)
|           |--------- Taxon2 (35,200)
|(30,50)
|---------------------- Taxon3 (20,100)
```

Then, the file with `.postmean1.tre` extension would encode the ancestral reconstruction as follows:

```
((Taxon1_40_150:0.4,Taxon2_35_200:0.4)56_80:0.6,Taxon3_20_100:1.0)_30_50;
```

and the file with the `.postmean1.tab` extension would read as:

```
Taxon1 Taxon1 0.4 40 150
Taxon2 Taxon2 0.4 35 200
Taxon1 Taxon2 0.6 56 80
Taxon3 Taxon3 1.0 20 100
Taxon1 Taxon3 0.0 30 50
```

Note that both files contain information about branch lengths. The tree is a time-calibrated phylogeny (in relative time), and the third column in the tabulated file corresponds to the length (in relative time) of the branch immediately upstream the last common ancestor of the two taxa (and 0 for the root).

The absolute time-calibrated phylogeny (with credibility intervals associated to internal nodes) is in the `postmeandates.tre` file.

We have also developed custom latex programs to draw chronograms, color trees and bubble trees out of these newick files (see our articles for examples). These are available upon request. However, compiling these programs, and compiling the resulting Latex files, requires special C++ and Tex libraries.

6 Detailed options of `coevol`

6.1 Input files and general settings

`-d <datafile>`

specifies a file containing aligned sequences (see: Input Data Format).

`-t <treefile>`

specifies the tree topology (see Input Data Format).

`-f`

forces the program to overwrite an already existing chain with same name.

`-x <every> [<until>]`

specifies the saving frequency, and (optional) the number of points after which the chain should stop. If this number is not specified, the chain runs "forever".

6.2 Substitution models

`-gc`

allows for variation in equilibrium gc frequencies of the substitution process between lineages.

`-tstv`

allows for variation in the ratio of transition over transversion rates between lineages.

`-dsom`

activates the codon model. The two a priori independent variables are dS and $\omega = dN/dS$ (see Lartillot and Poujol, 2011, for more details)

`-dsdn`

activates the codon model. The two a priori independent variables are dS and dN (see Lartillot and Poujol, 2011, for more details)

`-dsom2`

activates a codon model in which variation in dS and dN/dS are separately reconstructed for transitions and transversions.

`-dsom3`

activates a codon model in which variation in dS and dN/dS are separately reconstructed for transitions, GC-conservative and non GC-conservative transversions (see Lartillot, 2013). This model and the previous one (`-dsom2`) are computationally challenging and may require relatively large datasets to be correctly fitted.

`-mtvert`

selects the vertebrate mitochondrial genetic code

`-mtinv`

selects the 'invertebrate' mitochondrial genetic code

`-univ`

selects the universal genetic code

`-pol`

activates the amino-acid Kr/Kc (radical versus conservative) model. Amino-acid substitutions that do not conserve polarity are considered as radical (Nabholz et al., 2013).

`-vol`

activates the amino-acid Kr/Kc (radical versus conservative) model. Amino-acid substitutions that do not conserve volume are considered as radical (Nabholz et al., 2013).

`-polvol`

activates the amino-acid Kr/Kc (radical versus conservative) model. Substitutions that change the polarity and/or the volume of the amino-acid are considered as radical (Nabholz et al., 2013).

`-charge`

activates the amino-acid Kr/Kc (radical versus conservative) model. Substitutions that change the charge of the amino-acid are considered as radical (Nabholz et al., 2013).

6.3 Divergence times and branch lengths

`-cal <calibration_file>  <mean> <stdev>`

specifies the calibrations, jointly with the mean and standard deviation for the prior on the age of the root.

`-bd`

selects the birth death prior on divergence times.

`-unif`

selects the uniform prior on divergence times

6.4 Covariance matrix

`-diag`

constrains the covariance matrix to be a diagnoal matrix (all covariances between rates and traits equal to 0). Thus, all substitution parameters and quantitative characters now evolve along the tree in an uncorrelated fashion (for an example application, see Lartillot and Delsuc, 2012).

`-priorsigma <kappa>`

the prior on the covariance matrix is an Inverse Wishart distribution parameterized by a diagonal matrix Σ_0 and with q degrees of freedom. If $\kappa > 0$, $\Sigma_0 = \kappa I_M$. If $\kappa = -1$ (default option), each entry along the diagonal of the matrix Σ_0 is different, and all are estimated from the data (each being endowed with a truncated log-uniform prior between 10^{-3} and 10^3). This latter solution is better, as different traits or rates may have a very different rate of variation. Typically, log body size has a high rate of variation whereas $\ln \omega = dN/dS$ evolves within a very narrow range.

`-df <q>`

specifies the number of degrees of freedom q of the Inverse Wishart distribution (Lartillot and Poujol, 2011). In the current version, the default value is the number of entries of the multivariate process M. A properly defined Wishart distribution requires that $q > M - 1$.

6.5 Brownian process

`-arith`

uses the arithmetic averaging method (see Lartillot and Poujol, 2011, for more details) for computing branch-specific mean values of the substitution parameters.

`-geod`

uses the geodesic averaging method (see Lartillot and Poujol, 2011, for more details) for computing branch-specific mean values of the substitution parameters.

`-root <rootfile>`

specifies the prior for the values of the quantitative traits at the root of the tree. By default, the model assumes a truncated uniform prior over $[-100, 100]$ for the logarithm of the value of quantitative traits (and of substitution parameters) at the root (Lartillot and Poujol, 2011). However, in some cases, we may have independent information about ancestral traits. For instance, the distribution of the logarithm of body mass (in grams) of fossil mammals of the cretaceous in North America has a mean of 4.5 and a standard deviation of 1.5 (Alroy, 1999). In the context of our analysis of placental mammals above, this information could be used to derive a prior for the body mass at the root, that is, for the last common ancestor of placentals (although doing this amounts to ignoring the possibility that body-mass dependent diversification rates may result in this last common ancestor not being a random sample from cretaceous mammals, in terms of its body size).

Using the `-root` option, it is possible to enforce normal distributions for the log of the traits at the root, and to specify the means and variances of these normal distributions. The file should be specified as follows:

```
<Ntrait>
<mean1>   <stdev1>
<mean2>   <stdev2>
...
```

the means and standard deviations should be given trait by trait. Specifying a mean and a standard deviation of 0 will be equivalent to selecting a uniform distribution over $[-100, 100]$ for the corresponding trait.

7 Ancov: ancestral covariance and comparative regression

The `ancov` program takes as an input:

- a tree with branch lengths (that may or may not represent explicit time).
- a trait: a matrix of quantitative traits in extant taxa
- a predictor: a matrix of quantitative variables specified at both terminal and internal nodes of the phylogeny

The trait and the predictor can both be multivariate. The program then assumes that the joint evolution of the trait and the predictor is multivariate Brownian along the phylogeny. It relies on the joint knowledge of the trait and the predictor at the tips of the tree to estimate the covariance matrix of the process and the ancestral values of the trait along the phylogeny.

An example is given in the package. The `coevol/data/archaea/` directory contains the following files:

- `prot.tree`: a phylogenetic tree,
- `archaea.exptemp`: a matrix of exponentiated growth temperatures (remember that traits are log-transformed by default; here, however, we want to make a correlation with temperatures on a linear scale),
- `rna.itgc`: the inferred GC composition of rRNA stems along the tree. GC compositions (x) have been transformed as $y = \frac{x}{1-x}$,
- `prot.comp`: the inferred amino-acid composition of an amino-acid recoded multiple sequence alignment from Groussin and Gouy (2011).

As explained in Lartillot (2014), the GC and amino-acid compositions were inferred using the `phylobayes` program, under a non-homogeneous model of sequence evolution.

The following command:

```
ancov -t prot.tree -c archaea.exptemp -anc rrna.itgc rnatemp1
```

will run a MCMC under the model correlating the (logit-transformed) GC composition of rRNA stems with growth temperature (the latter, on a linear scale).

MCMC is much faster with `ancov` than with `coevol` (mostly because `ancov` does not involve any explicit likelihood computation under a model of sequence evolution, which is the rate-limiting aspect of any phylogenetic program). Thus, it is probably a good idea to subsample:

```
ancov -t prot.tree -c archaea.exptemp -anc rrna.itgc -x 10 -1 rnatemp2
```

Similar analyses can be conducted using protein composition. Here, however we have to be cautious concerning the exact parameterization. In Lartillot, 2014, the amino-acid compositions (such as given in `prot.comp`) were log-transformed and projected on the 19-dimensional hyperplane defined by $\sum_i x_i = 0$.

This transformation can be done automatically by `ancov`, using the `-ancfreq` (ancestral frequencies) command:

```
ancov -t prot.tree -c archaea.exptemp -anc prot.comp -ancfreq -x 10 -1 prottemp1
```

It works for any K-dimensional trait that sums to 1 over the K entries (thus with $K-1$ degrees of freedom, really).

Alternatively, the transformation can be done using a separate program, called `redcomp`, which is provided in the package. The command is:

```
redcomp <infile> <outfile>
```

so, here:

```
redcomp prot.comp prot.redcomp
```

will produce a `redcomp` file containing the 19-dimensional trait resulting from the log-transformation followed by the projection, which can then be used to run `ancov`:

```
ancov -t prot.tree -c archaea.exptemp -anc prot.redcomp -x 10 -1 prottemp1
```

This alternative way to proceed is less convenient that the automatic one in general, although it can be practical in specific situations, in particular if you want to combine several traits. For instance, in the present case, if you want to jointly reconstruct the evolution of temperature, ribosomal GC content and proteome composition along the archaeal tree, then the only reasonable approach is to first log-transform and reduce the proteome composition data, and then paste them with the rRNA GC data, thus producing a 20-dimensional trait that can then be used as an input to coevol.

All these analyses can be read using readancov:

```
readancov -x 100 -1 +log prottemp1
```

This command will produce a file named `prottemp1.postmean1.tre`, displaying the inferred ancestral growth temperatures along the tree. Here, a `+log` option is used, in order to obtain a reconstruction of temperatures (and not of their exponential).

Note that, in addition to the usual covariance and precision matrices, the .cov file produced by `ancov` will also contain an estimate of the proportion of the variance of each trait explained by all other traits. Thus, for instance, when correlating proteome composition and temperature:

```
proportion of variance of each trait explained by all other traits:
trait 0 : 0.85
trait 1 : 0.57
...
```

which means that no less than 85% of the variance of temperature among archaeal species is explained by proteome composition (and this is corrected for phylogenetic inertia).

8 Tipcoevol: a tip-dating version of coevol

The `tipcoevol` program implements the mixed relaxed clock model introduced in Lartillot et al. (2016), as well as a tip-dating version of `coevol`. In tip dating, fossils and extant taxa are all represented as tips of a *serial* phylogeny, endowed with a serial birth-death prior (with constant fossil sampling rate through time). Each fossil is constrained to be within a specified interval (through a separate calibration file). Molecular data (for extant taxa) and quantitative traits (for both extant and fossil taxa) are given as an input, and the program then infers divergence times as well as the correlation between substitution rate and quantitative traits (as in `coevol`). The program can be run without quantitative traits, in which case only a divergence estimation in a tip-dating context is done. With a specific option, the program also allows for node calibrations, exactly as in `coevol`. This allows using the mixed relaxed clock model in a node-dating context.

The program takes as an input:

- a serially sampled tree with branch lengths specified in absolute time (for tip-dating) or an ultrametric tree with only extant taxa (for node-dating).
- a calibration file specifying the intervals for the ages of the fossils (in the tip-dating case) or for the ages of the calibrated internal nodes (in the node-dating case). In the tip-dating case, the tree given as an input will be used as the starting tree. It should be compatible with the interval calibrations.
- a nucleotide sequence alignment (for tip-dating, this alignment should also contain the fossil taxa, with missing characters)
- (optional) a matrix of quantitative traits for extant and/or fossil taxa.

Examples (corresponding to the analysis done in Lartillot et al. (2016)) are provided in the `data/eutheria138` for the tip-dating case, and `data/eutheria105` for the node-dating case.

Detailed options

The following options have the same syntax and same meaning as in `coevol`:

- `-d`: for specifying a nucleotide sequence matrix
- `-c`: for specifying a dataset of quantitative traits
- `-t`: for speciyfing the phylogenetic tree
- `-cal`: for specifying the calibration file
- `-linear`, `-logit`: transforms the quantitative traits
- `-diag`: implements the diagonal model (without correlations between rate and traits)
- `-df`: gives the number of degreees of freedom of the inverse Wishart prior on the covariance matrix
- `-x`: specifies the burnin and subsampling frequency
- `-f`: forces the overriding of an already existing chain

Additional options, specific to `tipcoevol`:

`-ln`

implements the simple Brownian autocorrelated relaxed clock model (default option). Without quantitative traits, this model is exactly the log-normal autocorrelated clock of Thorne et al. (1998). With quantitative traits, this is the same model as in `coevol` program: a multivariate Brownian process whose first component corresponds to the substitution rate and the other components to each of the quantitative traits given as an input to the program.

`-lnwn`

implements the mixed relaxed clock model. If quantitative traits are given to the program, the first component of the multivariate Brownian process will map to the Brownian component of the mixed clock. This Brownian component will therefore be coupled with quantitative trait evolution, if quantitative traits are given as an input to the program.

`-wn <ratemean> <ratestdev>`

implements the pure white-noise model: the instant rate of substitution is a gamma white-noise process, of mean r_{wn} and standard deviation σ_{wn}. The prior on the standard deviation is an exponential of mean 1. As for the mean, its prior is a gamma distribution of mean and standard deviation specified by the used (`ratemean` and `ratestdev`). If quantitative traits are given to the program, the multivariate Brownian process will only describe their evolution (i.e. there will be no entry of this process mapping to the substitution rate).

`-nodebd <rootmean> <rootstdev> <bddivratemean> <bddivratestdev>`
`<mumean> <mustdev> <psimean> <psistdev> <rhomean> <rhostdev>`

activates the serial birth-death tip-dating model. The prior on the age of the root is a gamma distribution (of mean and standard deviation `rootmean` and `rootstdev`). The diversification rate of the birth-death prior ($r = \lambda - \mu$, difference between the speciation and the extinction rates), the extinction rate (μ), the sampling rate (ψ) and the sampling fraction at the present (ρ) are also endowed with a gamma prior (of mean and standard deviation given as arguments, such as specified above).

`-serialbd <rootmean> <rootstdev> <bddivratemean> <bddivratestdev> <mumean> <mustdev>`

activates the node-dating mode. The prior on the age of the root is a gamma distribution (of mean and standard deviation `rootmean` and `rootstdev`). The diversification rate of the birth-death prior ($r = \lambda - \mu$, difference between the speciation and the extinction rates) and the extinction rate (μ) are also endowed with a gamma prior (of mean and standard deviation given as arguments, such as specified above).

`-divsampling <cutofftime> <Ntaxa>`

activates the diversified sampling option. The total number of extant taxa of the group is estimated to be equal to `Ntaxa`. The extant taxa represented in the phylogeny given as an input to the program are assumed to be a representative set of all extant groups, all of which have coalesced onto one of the terminal branches of the specified phylogeny before the cutoff time. In other words, if n is the number of extant taxa represented in the phylogeny, the n ancestors obtained by following the n terminal branches backward in time until `cutofftime` are assumed to cover the whole of the extant diversity of the clade under study. Note that this implies that all internal nodes of the phylogeny should be older than `cutofftime`. The tree given as an input should be compatible with this requirement.

Reading chains produced by `tipcoevol`

Post-processing of the chains produced by `tipcoevol` is done using the `readtipcoevol` program. This program works like `readcoevol`. The only specific option is:

`-rate`

under the mixed clock model, this option will calculate the posterior mean proportion of rate variation explained by the Brownian and the white-noise components of the mixed clock. Note that, if most of the variance is contributed by the Brownian or by the white-noise component, this means that the program has effectively selected the relevant pure clock automatically. In other words, no explicit model comparison is necessary here, to select among alternative, autocorrelated or uncorrelated clocks.

9 Reconstructing variation in effective population size along phylogenies

This section introduces the models used in Brevet and Lartillot (2021), for reconstructing variation in effective population size (N_e) along a phylogeny, using data about synonymous polymorphism (π_S) and about generation time (τ) in extant species. The main idea behind this method is to rely on the following relation, implied by population genetics theory:

$$\pi_S \;=\; 4\,N_e\,u,$$

where u is the rate of mutation per generation. Assuming synonymous mutations are neutral, the mutation rate u can be estimated based on the synonymous substitution rate dS, combined with information about generation time. This simple idea has often been used to estimate N_e in a species of interest based on intra-specific polymorphism and (calibrated) synonymous divergence with a closely related species.

Here, the idea is just to do this in an integrative manner and in a way that allows all of these variables (N_e, u, dS, τ, etc) to vary continuously between species, according to a multivariate Brownian diffusion process. In practice, it is just a matter of giving as an input to `coevol` estimates of π_S and τ in extant species, along with other quantitative traits of interest. Then, the whole machinery of `coevol` is recruited to extract N_e and u from π_S based on the relations indicated above, while inferring the correlation patterns between N_e, u and all other quantitative variables of interest (dS, dN/dS, life-history traits etc) over the phylogeny.

As an example of how to conduct such an analysis, the data for the primate example explored in Brevet and Lartillot have been made available in the `coevol/data/primatesNe` subfolder, and the commands for reproducing this analysis are given below. The manuscript is available at `biorxiv.org/content/10.1101/793059v1`.

Two approaches were used in Brevet et al (2021):

- a phenomenological approach, in which the default multivariate Brownian model implemented by the `coevol` program is used, and then the `readcoevol` program is called with a special option (`-Ne`). This approach is directly available with the current implementation of `coevol`, such as available on the master branch.

- a mechanistic approach: here, a special model is implemented, in a separate program called `nearlyneutral`, with the post-analysis being implemented in `readnearlyneutral`. These programs are available on a separate git branch named `coevolNe`, which can be activated with the following command (when in the `coevol` folder):

```
git checkout coevolNe
```

and then recompiling the binaries by typing `make` when in the `sources` subfolder.

Of note, compared to the phenomenological approach, the mechanistic approach makes much stronger assumptions: essentially, a purely nearly-neutral regime (no positive selection) and a constant distribution of fitness effects of non-synonymous mutations (DFE) across species. As a result, it is more constrained and therefore gives more focussed ancestral N_e estimates (smaller credible intervals). On the other hand, it is potentially more sensitive to violations of these assumptions. This point is discussed in Brevet and Lartillot (2021). The phenomenological approach, on the other hand, is more agnostic about the underlying molecular evolutionary processes. It mostly assumes that synonymous mutations are neutral and that all variables of interest evolve according to a Brownian motion. Then, it is just using $\pi_S = 4N_e u$ to extract N_e and to infer its correlation patterns with other variables such as π_N/π_S and dN/dS, and this, without assuming a priori that these variables should follow any specific scaling relation with N_e. Thus, in practice, the phenomenological approach is probably the one to use by default.

9.1 Data

Both approaches take as an input:

- a multiple sequence alignment of coding sequences (codon compliant). For the primate example, we use a modified subset of the alignment published by Perelman et al. (2011).
- a rooted phylogenetic tree.
- data about synonymous and non-synonymous polymorphism in extant species (missing data are allowed). In the primate example, these data were obtained from Figuet et al. (2016). Specifically, the method requires a separate estimate of π_S and, optionally, an estimate of π_N/π_S.
- fossil calibrations: these are important for calibrating the synonymous substitution rate dS (taken as a proxy for the mutation rate) in terms of absolute time. Calibrations should be in My.
- data about generation times in extant species (in days): these will be used for converting synonymous substitution rates per My in mutation rates per generation (u). In turn, u will be used to get an estimate of N_e based on $\pi_S = 4N_e u$.

- possibly, other quantitative or life-history traits, which we might want to correlate with N_e. In the case of primates, data about generation times and life-history traits were obtained from the AnAge database (de Magalhaes and Costa, 2009).

The sequence alignment, the phylogenetic tree and the fossil calibrations should be formatted as for other analyses with `coevol` (see above, section 3, Input data format). Data about polymorphism, generation times and other quantitative traits should be provided as one single file, formatted as follows (column order does not matter, but the names for diversity and generation time are important):

```
#TRAITS
<Ntaxa> <Ntraits> trait_1 trait_2 ... trait_K piS piNpiS generation_time
Species_name <trait_1> .... <trait_K> <piS> <piNpiS> <generation_time>
...
```

`Ntraits` should be equal to the number of traits (including π_S, π_N/π_S and generation time), thus here `Ntraits` $= 6$. All traits and estimates of π_S and π_N/π_S should be given in natural units (they will be log-transformed by the program).

9.2 Phenomenological approach (`coevol` and `readcoevol`)

This approach simply runs the standard model of `coevol`. The output is then used to estimate N_e and its correlation patterns with other traits. The main idea is to rely on the following relation, implied by population genetics theory:

$$\pi_S \;=\; 4\,N_e\,u, \tag{1}$$

where u is the rate of mutation per generation. Assuming synonymous mutations are neutral, u can be estimated as:

$$u \;=\; dS\,\tau. \tag{2}$$

Finally, taking the logarithm:

$$\ln N_e \;=\; \ln \pi_S - \ln d_S - \ln \tau - \ln 4. \tag{3}$$

gives $\ln N_e$ as a linear combination of the components of the Brownian process, which can thus be reconstructed by simply making a log-linear change of variables (see Brevet and Lartillot, 2021, for more details).

9.2.1 Running the analysis

To run the analysis with the primate example, use the command:

```
coevol -d prim.phy -t prim.rootedtree -c prim.lht -cal prim.calib 80 80 -dsom primne1
coevol -d prim.phy -t prim.rootedtree -c prim.lht -cal prim.calib 80 80 -dsom primne2
```

Here, two independent chains are run, with names `primne1` `primne2`. In the analysis presented in Brevet and Lartillot, runs were conducted for 4000 cycles (run of about one week). However, a run of one day will give 600 cycles, which, with a burnin of 100 and a sample size of 500, will already give reasonably accurate estimates.

The post-analysis is done using `readcoveol` with the `-Ne` option:

```
readcoevol -x 100 1 -Ne primne1
readcoevol -x 100 1 -Ne primne2
```

Here, a burnin of 100 is used. This will output a tree with posterior credible intervals for N_e attached to all nodes (in the file named `primne1.postmeanNe.tre`), as well as a correlation analysis of N_e and u with other variables (in `primne1.cov`). The options described in section 5.4 can be used to modulate the output (mean, median, credible intervals, etc).

9.3 Mechanistic approach (`nearlyneutral` and `readnearlyneutral`)

This alternative approach relies on an explicit model of the nearly-neutral evolutionary process. It uses the original Coevol framework but introduces additional constraints, such that some of the parameters are deduced through deterministic relations implying other Brownian dependent parameters.

Specifically, the Brownian free variables are now:

$$\begin{aligned} X_1 &= \ln \pi_S \\ X_2 &= \ln \pi_N/\pi_S \\ X_3 &= \ln \tau \end{aligned}$$

Then, we make use a theoretical result (Kimura, 1979; Welch et al., 2008), showing that, assuming a purely nearly-neutral regime and a constant gamma-shaped distribution of fitness effects (DFE), π_N/π_S and d_N/d_S both scale allometrically as a function of N_e, with a scaling exponent equal to the shape of the DFE. The shape parameter is classically noted β. Thus:

$$\begin{aligned} \pi_N/\pi_S &= \kappa_2 N_e^{-\beta}, \\ dN/dS &= \kappa_1 N_e^{-\beta}, \end{aligned}$$

where κ_1 and κ_2 are two empirical constants. Using these relations, the other variables of interest can be expressed as deterministic functions of the three free variables defined by X, as:

$$\begin{aligned} \ln N_e &= -1/\beta \left(\ln \pi_N/\pi_S + \ln \kappa_2 \right), \\ \ln dS &= \ln \pi_S - \ln 4N_e - \ln \tau, \\ \ln dN/dS &= -\beta \ln N_e + \ln \kappa_1. \end{aligned}$$

The model has thus three structural parameters, β, κ_1 and κ_2, representing the structure of the DFE (assumed constant across the entire clade). Of note, in principle, there is a numerical constraint between κ_1 and κ_2 (Welch et al., 2008). However, if dS and dN/dS, on the one hand, and π_S and π_N/π_S, on the other hand, are not estimated on the same genes, these two constants have no reason to fulfill this constraint. Therefore, in the model introduced here, they are estimated separately. The three parameters β, κ_1 and κ_2 all have a normal prior, of mean 0 and variance 1.

9.3.1 Running the analysis

To run the analysis with the primate example, use the following command:

```
nearlyneutral -d prim.phy -t prim.rootedtree -c prim.lht -cal prim.calib 80 80 primne_mech1
nearlyneutral -d prim.phy -t prim.rootedtree -c prim.lht -cal prim.calib 80 80 primne_mech2
```

Once these chains have run for sufficiently long (600 to 1100 cycles), the post-analysis is done using `readnearlyneutral`:

```
readnearlyneutral -x 100 1 primne_mech1
readnearlyneutral -x 100 1 primne_mech2
```

This will output a tree with posterior credible intervals for N_e attached to all nodes (in the file named `primne_mech1.postmeanNe.tre`). The options described in section 5.4 can be used to modulate the output (mean, median, credible intervals, etc).

References

Alroy, J. 1999. The fossil record of North American mammals: evidence for a Paleocene evolutionary radiation. *Syst. Biol.* 48:107–118.

Dempster, AP. 1972. Covariance selection. *Biometrics* 28:157–175.

Diaz-Uriarte, R, and T Garland. 1998. Effects of branch length errors on the performance of phylogenetically independent contrasts. *Syst. Biol.* 47:654–672.

Drummond, Alexei J, and Andrew Rambaut. 2007. BEAST: Bayesian evolutionary analysis by sampling trees. *BMC Evol. Biol.* 7:214.

Drummond, D, Alpan Raval, and Claus Wilke. 2006. A single determinant dominates the rate of yeast protein evolution. *Mol. Biol. Evol.* 23:327–337.

Felsenstein, J. 1981. Evolutionary trees from DNA sequences: a maximum likelihood approach. *J. Mol. Evol.* 17:368–376.

Felsenstein, J. 1985. Phylogenies and the comparative method. *Am Nat* 125:1–15.

Figuet, Emeric, Benoit Nabholz, Manon Bonneau, Eduard Mas Carrio, Krystyna Nadachowska-Brzyska, Hans Ellegren, and Nicolas Galtier. 2016. Life History Traits, Protein Evolution, and the Nearly Neutral Theory in Amniotes. *Mol. Biol. Evol.* .

Franks, Peter J, Rob P Freckleton, Jeremy M Beaulieu, Ilia J Leitch, and David J Beerling. 2012. Megacycles of atmospheric carbon dioxide concentration correlate with fossil plant genome size. *Philos. Trans. R. Soc. Lond. B Biol. Sci.* 367:556–564.

Geyer, C. 1992. Practical Markov Chain Monte Carlo. *Stat. Sci.* 7:473–483.

Groussin, Mathieu, and Manolo Gouy. 2011. Adaptation to environmental temperature is a major determinant of molecular evolutionary rates in archaea. *Mol. Biol. Evol.* 28:2661–2674.

Harvey, P, and Mark Pagel. 1991. *The comparative method in evolutionary biology*.

Heath, Tracy A, John P. Huelsenbeck, and Tanja Stadler. 2014. The fossilized birth-death process for coherent calibration of divergence-time estimates. *Proc. Natl. Acad. Sci. USA* 111:E2957–66.

Kimura, M. 1979. Model of effectively neutral mutations in which selective constraint is incorporated. *Proc. Natl. Acad. Sci. USA* 76:3440–3444.

Lartillot, Nicolas. 2013. Interaction between Selection and Biased Gene Conversion in Mammalian Protein-Coding Sequence Evolution Revealed by a Phylogenetic Covariance Analysis. *Mol. Biol. Evol.* 30:356–368.

Lartillot, Nicolas. 2014. A phylogenetic Kalman filter for ancestral trait reconstruction using molecular data. *Bioinformatics* .

Lartillot, Nicolas, and Frédéric Delsuc. 2012. Joint reconstruction of divergence times and life-history evolution in placental mammals using a phylogenetic covariance model. *Evolution* 66:1773–1787.

Lartillot, Nicolas, Thomas Lepage, and Samuel Blanquart. 2009. PhyloBayes 3: a Bayesian software package for phylogenetic reconstruction and molecular dating. *Bioinformatics* 25:2286–2288.

Lartillot, Nicolas, Matthew J Phillips, and Fredrik Ronquist. 2016. A mixed relaxed clock model. *Philos. Trans. R. Soc. Lond. B Biol. Sci.* 371.

Lartillot, Nicolas, and Raphaël Poujol. 2011. A phylogenetic model for investigating correlated evolution of substitution rates and continuous phenotypic characters. *Mol. Biol. Evol.* 28:729–744.

de Magalhaes, J, and J Costa. 2009. A database of vertebrate longevity records and their relation to other life-history traits. *J. Evol. Biol.* 22:1770–1774.

Martins, E, and T Hansen. 1997. Phylogenies and the comparative method: a general approach to incorporating phylogenetic information into the analysis of interspecific data. *Am Nat* 149:646–667.

Nabholz, Benoit, Nicole Uwimana, and Nicolas Lartillot. 2013. Reconstructing the phylogenetic history of long-term effective population size and life-history traits using patterns of amino acid replacement in mitochondrial genomes of mammals and birds. *Genome Biol. Evol.* 5:1273–1290.

Organ, Chris L, Andrew M Shedlock, Andrew Meade, Mark Pagel, and Scott V Edwards. 2007. Origin of avian genome size and structure in non-avian dinosaurs. *Nature* 446:180–184.

Pagel, M. 1999. Inferring the historical patterns of biological evolution. *Nature* 401:877–884.

Perelman, Polina, Warren E Johnson, Christian Roos, Hector N Seuánez, Julie E Horvath, Miguel A M Moreira, Bailey Kessing, Joan Pontius, Melody Roelke, Yves Rumpler, Maria Paula C Schneider, Artur Silva, Stephen J O'Brien, and Jill Pecon-Slattery. 2011. A Molecular Phylogeny of Living Primates. *PLoS Genet.* 7:e1001342.

Pyron, R. 2010. A likelihood method for assessing molecular divergence time estimates and the placement of fossil calibrations. *Syst. Biol.* 59:185–194.

Ronquist, Fredrik, Seraina Klopfstein, Lars Vilhelmsen, Susanne Schulmeister, Debra L Murray, and Alexandr P Rasnitsyn. 2012. A total-evidence approach to dating with fossils, applied to the early radiation of the hymenoptera. *Syst. Biol.* 61:973–999.

Stadler, T, and Z Yang. 2013. Dating Phylogenies with Sequentially Sampled Tips. *Syst. Biol.* 62:674–688.

Stadler, Tanja. 2010. Sampling-through-time in birth–death trees. *J Theor Biol* .

Thorne, J L, H Kishino, and I S Painter. 1998. Estimating the rate of evolution of the rate of molecular evolution. *Mol. Biol. Evol.* 15:1647–1657.

Thorne, Jeffrey L, and Hirohisa Kishino. 2002. Divergence time and evolutionary rate estimation with multilocus data. *Syst. Biol.* 51:689–702.

Welch, John J, Adam Eyre-Walker, and David Waxman. 2008. Divergence and polymorphism under the nearly neutral theory of molecular evolution. *J. Mol. Evol.* 67:418–426.

Wong, Frederick, Christopher K Carter, and Robert Kohn. 2003. Efficient estimation of covariance selection models 90:809.

Yang, Ziheng, and B Rannala. 2006. Bayesian estimation of species divergence times under a molecular clock using multiple fossil calibrations with soft bounds. *Mol. Biol. Evol.* 23:212–226.